


Topic 1: Introduction

Course map, and mathematical foundations

Digital Signal Processing for Sound and Music · Topic 1

Xavier Serra · Universitat Pompeu Fabra

What is audio signal processing?




Intentional alteration, analysis, representation, or generation of sound signals.


Analog and digital signals

Analog signals vary continuously in time and amplitude; digital audio represents them as a sequence of numbers.

Continuous waveform



Sampled waveform



- Sampling rate fixes time resolution: e.g., 44.1 kHz means 44,100 samples per second.
- Quantization fixes amplitude resolution: each sample is stored with finite precision.

Core application areas

Storage

Represent sound files and streams reliably.

Compression

Use signal structure and perception to reduce data.

Effects

Alter timbre, space, pitch, or timing.

Synthesis

Generate sound from models and parameters.

Description

Extract features for search, MIR, and machine learning.

Course structure: 10 topics

1 Intro + math

2 DFT

**3 Fourier
properties**

4 STFT

**5 Sinusoidal
model**

**6 Harmonic
model**

**7 Sinusoidal +
residual**


8 Transformations

**9 Sound/music
description**

**10 Concluding
topics**


The sequence moves from mathematical representation to analysis/synthesis models, transformations, and descriptions.

Learning workflow every week



- Start from a signal-processing idea.
- Implement a minimal example.
- Inspect time-domain and spectral plots.
- Listen to the result and interpret what changed.

SMSTools: the course software ecosystem



The package provides sound analysis/synthesis tools in Python; the materials repository provides notebooks, exercises, sounds, and GUIs for using them in class.

Course expectation: use the tools, but understand the signal flow, parameters, plots, and listening consequences of every step.

What is inside sms-tools?

Fourier foundations

dftModel.py
stft.py
spectrum and sound
examples

Sound models

sineModel.py
harmonicModel.py
stochasticModel.py

Combined models

sinusoidal + residual
sinusoidal + stochastic
harmonic +
residual/stochastic


Transformations

analysis → modification →
synthesis
pitch, time, timbre
examples
GUI backends as
compact scripts

The names of the modules are not just file names: they mirror the conceptual sequence of the course, from DFT and STFT to sinusoidal, harmonic, residual, stochastic, and transformation models.

When a notebook calls a model, ask: What representation is being computed? Which parameters control it? What changes when I listen?

sms-tools-materials: where exercises happen



Typical setup


```
python -m venv .venv  
pip install -r requirements.txt  
jupyter notebook
```

Typical use in class

Open the exercise notebook, run the known example first, inspect the plots, listen to the output, then modify one parameter at a time.

Use your own audio only after the method works on the supplied examples.

How to use SMSTools in an exercise



Before changing code

Identify input $x[n]$ and sampling rate
Find the analysis parameters
Understand the plotted quantities

After changing code

Compare plots before/after
Listen to input and output
State the audible consequence

A successful exercise is not only “the cell runs”: it means you can connect the code to the signal-processing concept and the sound you hear.

Mathematics: why these tools?

The DFT, STFT, and spectral models are built from a compact set of mathematical ideas.

sinusoids

complex numbers

Euler formula

complex sinusoids

dot products

orthogonality

**even/odd
symmetry**

convolution

Sinusoidal functions

Continuous-time signal

$$x(t) = A \cos(2\pi f_0 t + \phi)$$

Discrete-time signal


$$x[n] = A \cos(2\pi f_0 n / f_s + \phi)$$

A: amplitude

f_0 : frequency in Hz

ϕ : phase shift

f_s : sampling rate, which maps time into sample index n



A sinusoid is fully described by amplitude, frequency, and phase. In digital audio we work with the sampled version $x[n]$.


Complex numbers and the complex plane

$$x = a + jb$$

$$\text{Re}\{x\} = a, \quad \text{Im}\{x\} = b, \quad j^2 = -1$$

A complex number is represented as a point or vector on the complex plane.

The horizontal axis is real; the vertical axis is imaginary.



Rectangular and polar forms

Rectangular

$$x = a + jb$$


Good for addition and subtraction: real and imaginary parts are explicit.

Polar

$$x = Ae^{j\phi}$$

$$A = |x| = \sqrt{a^2 + b^2} \quad \phi = \arg(x) = \text{atan2}(b, a)$$

Good for magnitude, phase, multiplication, and rotation. This is the form used to describe spectral magnitude and phase.




Euler's formula: rotation and sinusoids

$$e^{j\phi} = \cos(\phi) + j\sin(\phi)$$

$$\cos \phi = \frac{e^{j\phi} + e^{-j\phi}}{2}$$

$$\sin \phi = \frac{e^{j\phi} - e^{-j\phi}}{2j}$$

A complex exponential is a rotating point. Its real and imaginary projections are cosine and sine.



Real and complex sinusoids

Complex sinusoid

$$\bar{x}[n] = Ae^{j(\omega nT + \phi)} = Ae^{j\phi} e^{j\omega nT} = Xe^{j\omega nT}$$


$$\bar{x}[n] = A\cos(\omega nT + \phi) + jA\sin(\omega nT + \phi)$$

Real sinusoid


$$x[n] = A\cos(\omega nT + \phi) = A\left(\frac{e^{j(\omega nT + \phi)} + e^{-j(\omega nT + \phi)}}{2}\right)$$

$$= \frac{1}{2}Xe^{j\omega nT} + \frac{1}{2}X^*e^{-j\omega nT} = \frac{1}{2}\bar{x}[n] + \frac{1}{2}\bar{x}^*[n] = \Re\{\bar{x}[n]\}$$

Rotating complex exponential



Real and imaginary projections



The real signal is the real part of the complex sinusoid; the imaginary part is the quadrature projection.

Scalar product of sequences

Definition

$$\langle x, y \rangle = \sum x[n] y^*[n]$$

sum over $n = 0, \dots, N-1$

The scalar product measures similarity between two sequences.

Example with complex-valued sequences

$$x[n] = [0, j, 1] \quad y[n] = [1, j, j]$$

$$y^*[n] = [1, -j, -j]$$

$$\langle x, y \rangle = 0 \cdot 1 + j \cdot (-j) + 1 \cdot (-j) = 1 - j$$

For complex sequences, the second sequence is conjugated before multiplication.

Element-by-element products

n	0	1	2
x[n]	0	j	1
y*[n]	1	-j	-j
x[n] y*[n]	0	1	-j

$$\text{sum} = 0 + 1 - j = 1 - j$$

Orthogonality of sequences

Definition

$$\mathbf{x} \perp \mathbf{y} \Leftrightarrow \langle \mathbf{x}, \mathbf{y} \rangle = 0$$

Orthogonal sequences have zero scalar product.


Simple geometric example

$$\mathbf{x} = (2, 2) \quad \mathbf{y} = (2, -2)$$

$$\langle \mathbf{x}, \mathbf{y} \rangle = 2 \cdot 2 + 2 \cdot (-2) = 0$$

The two vectors are perpendicular. This is the geometric intuition behind orthogonal basis functions in the DFT.

Vector view of the example



Because $\langle \mathbf{x}, \mathbf{y} \rangle = 0$, the two directions do not “share” any component.

Even and odd functions

Definitions


Even: $x[-n] = x[n]$

Odd: $x[-n] = -x[n]$

Decomposition


$$x_e[n] = \frac{x[n] + x[-n]}{2} \quad x_o[n] = \frac{x[n] - x[-n]}{2}$$

Any signal can be decomposed into an even part and an odd part. These symmetries are useful when interpreting Fourier-transform properties.



Convolution in discrete time


$$y[n] = (x_1[n] * x_2[n])_n = \sum_{m=0}^{N-1} x_1[m]x_2[n-m]$$



Convolution combines an input signal with a system response.

In LTI systems, it gives the output for any input.

The first big conceptual bridge



Fourier analysis is a change of viewpoint: from samples over time to weights of sinusoidal building blocks.

Recommended references for Topic 1

- MTG sms-tools package: Python sound analysis/synthesis tools used throughout the course.
- MTG sms-tools-materials repository: lecture material, exercises, example notebooks, sounds, and GUI interfaces.
- Julius O. Smith, Mathematics of the DFT: a rigorous online reference for DFT foundations.
- Wikipedia background refreshers: audio signal processing, sinusoid, complex numbers, Euler formula, dot product, convolution.
- Audacity and sound examples can be used to listen to basic transformations and effects.

Exercise 1: Python and sounds

First lab: move from formulas to samples, code, plots, and listening.

What you will practice

- read a mono WAV file into a NumPy array
- inspect samples, indexing, amplitude range, and plots
- write small functions for reproducible audio operations
- downsample sounds and listen to the consequences

How to start

Open: `exercises/E1-Python-and-sounds.ipynb`

Use: NumPy · Matplotlib · IPython Audio · `wavread()`

Sounds: piano.wav, oboe-A4.wav, vibraphone-C6.wav, sawtooth-440.wav

Lab structure

1 · Reading audio

- `read_audio()`
- return sample slices

2 · Basic audio ops

- `min_max_audio()`
- plot and inspect waveform

3 · Array indexing

- `hop_samples(x, M)`
- select every Mth element

4 · Downsampling

- $M = 14$
- listen for aliasing

Listening question: what changes when we keep fewer samples, and why do different sounds react differently?