

The Short-Time Fourier Transform

Topic 4 · Analysis windows, spectrograms, inverse STFT, FFT size and hop size

Digital Signal Processing for Sound and Music · Topic 4

Xavier Serra · Universitat Pompeu Fabra

STFT

Lecture roadmap

1 STFT equation

local spectra from overlapping frames

2 Analysis windows

window shape and spectral leakage

3 M, N, H

window size, FFT size, and hop size

4 Spectrograms

time-frequency compromise, magnitude and phase

5 Inverse STFT

frame synthesis and overlap-add

6 Example + Exercise 4

stft-examples.ipynb and E4-STFT

Goal: understand the STFT as repeated windowed DFT analysis, and learn how window, FFT size and hop size change what we see, hear, and reconstruct.

Short-Time Fourier Transform equation

$$X_{\ell}[k] = \sum_{n=-M/2}^{M/2-1} x[n + \ell H] w[n] e^{-j2\pi kn/N}$$

$k = -N/2, \dots, N/2 - 1$ $\ell = 0, 1, \dots$

The STFT is not a new transform: it is a sequence of DFTs, each applied to a short windowed frame of the signal.

Symbols

$w[n]$: analysis window

ℓ : frame index

H : hop size

M : window length; N : FFT size

Interpretation

For each frame ℓ , take a local segment centered around time ℓH , multiply it by the window, then compute one DFT column.

Phase convention

Centered indexing around $n = 0$ makes phase plots easier to interpret and avoids artificial phase ramps.


Why short-time analysis?

Most musical sounds are not stationary: their spectrum changes over time.

The STFT assumes the signal is locally stable inside a short frame.

Frames overlap so changes can be tracked smoothly.

This is the conceptual bridge from the DFT to spectrograms: one DFT is one frame; an STFT is a sequence of DFTs.



Transform of a windowed sinewave

$$\begin{aligned}x_w[n] &= w[n]A_0\cos(2\pi k_0 n/N) \\&= \frac{A_0}{2}w[n]e^{j2\pi k_0 n/N} + \frac{A_0}{2}w[n]e^{-j2\pi k_0 n/N}\end{aligned}$$

$$X_w[k] = \frac{A_0}{2}W[k - k_0] + \frac{A_0}{2}W[k + k_0]$$

The phrase “spectrum of a frame” always means spectrum of the windowed frame, not of the raw signal segment alone.

Main idea

A sinusoid multiplied by a window does not produce an ideal impulse in frequency.


What appears in frequency

The spectrum is two shifted copies of the window spectrum $W[k]$, one at $+k_0$ and one at $-k_0$.

Consequence

The window controls leakage and peak shape: main-lobe width and side-lobe level.

Windowed sinewave: what the formula looks like




The time signal is a sinusoid tapered by the analysis window.

The magnitude spectrum has peaks around $\pm k_0$, shaped by $W[k]$.

The phase is meaningful mainly where magnitude is significant.

This is why magnitude-only plots can hide important assumptions about windowing and phase.

Analysis window



Window shape

A window reduces discontinuities at frame boundaries, but it also has its own spectrum.

Main lobe

Controls frequency resolution: narrow lobes separate close sinusoids better.

Side lobes

Control leakage: lower side lobes reduce interference from stronger nearby components.

Window functions in SciPy

Common `scipy.signal` window functions used in STFT analysis

boxcar / rectangular	maximum resolution, high leakage	flattop	amplitude accuracy
bartlett	triangular window	gaussian	Gaussian shape, parameter std
barthann	modified Bartlett-Hann	general_gaussian	generalized Gaussian shape
hann / hanning	good general-purpose taper	kaiser	tunable beta parameter
hamming	lower first side-lobe	chebwin	tunable attenuation
blackman	stronger side-lobe suppression	parzen	Parzen window
blackmanharris	very low side-lobes	triang	triangular window
bohman	Bohman window	slepian / dpss	energy concentration
nuttall	minimum 4-term Blackman-Harris		

In code: `scipy.signal.get_window(name, M)`. Always report the window name together with M, N, and H.

Rectangular window


$$w[n] = 1, \quad n = -M/2, \dots, M/2$$

$$W[k] = \frac{\sin(\pi k)}{\sin(\pi k/M)}$$

main-lobe width: 2 bins

side-lobe level: -13.3 dB

Best frequency resolution but poor leakage control.



Hann window


$$w[n] = 0.5 + 0.5\cos(2\pi n/M)$$

$$W[k] = 0.5D[k] + 0.25(D[k-1] + D[k+1])$$

main-lobe width: 4 bins

side-lobe level: -31.5 dB

A common default for spectrograms.



Hamming window


$$w[n] = 0.54 + 0.46\cos(2\pi n/M)$$

$$W[k] = 0.54D[k] + 0.23(D[k-1] + D[k+1])$$

main-lobe width: 4 bins

side-lobe level: -42.7 dB

Lower first side-lobe than Hann.




Blackman window

$$w[n] = 0.42 - 0.5\cos(2\pi n/M) + 0.08\cos(4\pi n/M)$$

main-lobe width: 6 bins

side-lobe level: -58 dB

Good leakage suppression, wider peaks.



Blackman-Harris window

$$w[n] = a_0 - a_1 \cos\left(\frac{2\pi n}{M-1}\right) + a_2 \cos\left(\frac{4\pi n}{M-1}\right) - a_3 \cos\left(\frac{6\pi n}{M-1}\right)$$

$$a_0 = 0.35875, a_1 = 0.48829, a_2 = 0.14128, a_3 = 0.01168, \quad 0 \leq n < M$$


main-lobe width: 8 bins

side-lobe level: about -92 dB


Excellent side-lobe suppression, but lower frequency resolution.

Interpretation

Use it when weak components near strong components matter more than resolving very close frequencies.



Window choice changes the spectrum



The same frame produces different spectra depending on the window.


Rectangular windows show strong side-lobes.

Hamming and Blackman reduce leakage but widen peaks.

No window is best for every task.

Always interpret a spectral plot together with the window that produced it.

STFT and analysis window



Local analysis

The STFT converts each short windowed fragment into a complex spectrum.


What changes from frame to frame?

The frame center IH moves through the sound; each new frame produces a new column in the spectrogram.

What remains fixed?

The analysis parameters: window type, M , N , and H .

Window size




Small M : better time localization, poorer frequency discrimination.

Large M : better frequency discrimination, poorer time localization.

The “right” M depends on how fast the sound changes and which frequencies matter.

This is the central time-frequency compromise of the STFT.

Even-odd size window




The window length M can be even or odd.

Centering conventions affect how samples are arranged around $n = 0$.

For phase-sensitive analysis, be explicit about indexing and zero-phase placement.

In implementation, a one-sample mismatch can introduce a visible phase difference.

FFT size



N controls the number of DFT samples used to represent the spectrum.

When $N > M$, the frame is zero-padded before the FFT.

Zero-padding interpolates the spectrum; it does not add new information.

Rule of thumb

M controls analysis-window duration; N controls how densely the spectrum is sampled.


Hop size

$$A_w[n] = \sum_{l=0}^{L-1} w[n - lH] = c$$


H is the distance in samples between consecutive frame centers.

Smaller H means more overlap and smoother time tracking.

For good resynthesis, overlapped windows should sum to a near-constant value.



Time-frequency compromise




Short windows follow transients and rapid changes.

Long windows distinguish close frequencies more clearly.

The spectrogram is never “the signal”; it is a view determined by M, N, H, and the window.

When comparing spectrograms, always compare analysis parameters first.

Amplitude and phase spectrograms



Magnitude spectrograms show energy distribution over time and frequency.

Phase spectrograms store temporal alignment information.

Phase often looks noisy when plotted directly, but it is essential for reconstruction and transformations.

Do not interpret phase where magnitude is close to zero.

Inverse STFT

$$\hat{x}_\ell[n] = \frac{1}{N} \sum_{k=-N/2}^{N/2-1} X_\ell[k] e^{j2\pi kn/N}$$

$$\tilde{x}_\ell[n] = \hat{x}_\ell[n - \ell H] w_s[n - \ell H]$$

$$y[n] = \frac{\sum_{\ell} \tilde{x}_\ell[n]}{\sum_{\ell} w[n - \ell H] w_s[n - \ell H]}$$

Implementation logic: IFFT each frame → place it in time → overlap-add → compensate the window sum.

1 • Inverse DFT

Convert each spectral column $X_\ell[k]$ back into a time-domain frame.

2 • Window + shift


Apply the synthesis window, move the frame to time ℓH , and add it to the output buffer.

3 • Normalize

Divide by the overlap of analysis/synthesis windows. With the same window this becomes a sum of w^2 terms.

Analysis/synthesis with magnitude and phase

$$X_l[k] \rightarrow (mX_l[k], pX_l[k]) \rightarrow y[n]$$




The complex STFT contains magnitude and phase for every frame.

IFFT reconstructs each frame.

Overlap-add reconstructs the full time signal.

Magnitude-only reconstruction is generally not exact.


STFT system



Analysis creates spectra on a frame grid; synthesis uses the same grid. Exact reconstruction needs phase plus proper overlap-add normalization.

Example notebook: STFT workflow

examples/stft-examples.ipynb connects the STFT equations to a complete experiment: analyze, inspect, synthesize, and listen.



Reading strategy: every plot should be tied to the exact windowed frame, centering convention, FFT size, and hop size that produced it.

How to read the STFT example outputs

1 • Windowing

The spectrum belongs to the windowed frame. Changing the window changes leakage and peak shape.

2 • Frame and hop

M determines how much sound each spectrum sees; H determines how often spectra are computed.

3 • Magnitude and phase

Magnitude makes the time-frequency structure visible; phase is required for reconstruction.

4 • Symmetry

For real input frames, positive and negative frequencies are conjugate-symmetric before one-sided display.

5 • Resynthesis

The IFFT and overlap-add steps test whether the analysis representation preserves the signal.

6 • Parameter sensitivity

Change M, N, H and window type one at a time; listen and inspect the result.

What to remember from Topic 4

1 STFT

a sequence of windowed DFTs

2 Window

controls leakage and time-domain framing

3 M

time/frequency analysis window length

4 N

FFT size and spectral sampling density

5 H

hop size and overlap-add behavior

6 ISTFT

reconstruction needs magnitude, phase, and compatible windows

The STFT is an analysis method, not a neutral picture: every spectrogram reflects the chosen frame, window, FFT size and hop size.

Exercise 4 — Short-Time Fourier Transform

Notebook: [exercises/E4-STFT.ipynb](#)

Goal: turn the STFT equations into code and learn how M , N , H and the window shape affect analysis and reconstruction.

window
 M

FFT
 N

hop
 H

1 · STFT analysis

Compute windowed frames and spectra for a sound.

2 · Spectrograms

Plot magnitude and phase over time.

3 · ISTFT

Synthesize frames and overlap-add them.


4 · Parameters

Compare M , N , H , and window choices.

Main check: for every plotted spectrum, identify the exact frame, window, FFT size and hop size that produced it.

What you will implement / verify in Exercise 4

- Use a window function and frame index to extract $x[n+lH]$.
- Compute one DFT per frame and store magnitude and phase.
- Generate amplitude and phase spectrograms.
- Reconstruct the signal using inverse transforms and overlap-add.
- Compare reconstruction error and listen to the output.
- Explain how M , N , H and window type changed the plots.



Focus on controlled experiments: change one parameter at a time, then inspect and listen.

References and credits

- Original MTG sms-tools-materials lecture slides and plotting code.
- Example notebook: `examples/stft-examples.ipynb`
- Exercise notebook: `exercises/E4-STFT.ipynb`
- Wikipedia background: STFT, window function, spectrogram.
- Julius O. Smith, Spectral Audio Signal Processing: STFT and windowing chapters.
- Sounds from the sms-tools-materials sounds folder and Freesound pack references.
- Slides follow the CC BY-NC-SA spirit of the original materials; code references follow AGPL licensing.

Keep references practical: use them when an implementation detail or a plot interpretation is unclear.